

Adoption playbook — putting SSSF to work on real code


Adoption playbook

- Rule zero: wire the gates before an agent writes a line
- Part A — an existing codebase
 - A1. Recon before spending
 - A2. Wire the gates (rule zero)
 - A3. Set the boundaries
 - A4. Interrogate before you plan
 - A5. The four jobs
 - A6. Read what happened
- Part B — a new project
 - B1. Think before there is code (interactive, not SSSF)
 - B2. Make it a repo with a walking skeleton
 - B3. Install and wire (rule zero)
 - B4. Tune the roster
 - B5. Run the slices, one at a time
- The standing checklist
- Where the two layers sit
- Part C — going dark: bootstrap vocabulary once, then let the loop run headless
 - Problem 1 — the interview a vendored planning skill wants doesn't have anyone to answer it
 - Problem 2 — nothing stops a plan from re-implementing what already exists
 - The grill → triage handoff is manual, on purpose (accepted, not a gap to close)
 - Definition of done, extended
- Part D — the queue: from `ready-for-agent` to shipped
 - The whole chain, named once
 - Filing: less manual than it looks, but one real gap remains
 - Two owners of one file — protect the queue's extensions on re-sync

- Mandatory checkpoints — verify the filing actually happened
- The judgment call stays interactive, on purpose
- The queue mechanism already exists — reuse it
- What `just watch` does, concretely
- Definition of done, extended again
- What this playbook does not claim
- Cheat sheet — the five vendored skills
- Cheat sheet — the extended Pocock+ToB roster (validated downstream)
- Version history

Adoption playbook

Two situations: a codebase that already exists, and a project that does not yet. Both start from the same non-negotiable rule.



Rule zero sits early in both parts, not at the end: gates are wired before any agent writes a line.

Rule zero: wire the gates before an agent writes a line

A freshly stamped factory ships every quality block as an `echo` that exits 0. Run a workflow against it and you get a green trace that proves nothing — the agent's claim went unchecked and the run *reported success*. This is worse than having no factory, because it manufactures confidence.


So the first thing you do in any repo, before any agent writes anything:

```
just quality "baseline"           # zero agents – look at what it claims
```

If it says `PLACEHOLDER`, the gates are fake. Wire them in `adws/adw_modules/quality.py`, then **prove each one can fail**:

```
# break exactly one thing, run the gates, confirm exactly one goes red,
  restore
```

A gate you have not watched fail is not a gate. If a gate stays green against a deliberately broken build, treat the *check* as the defect, not the tool: the check does not reach that class of breakage. Do this step before anything else in the adoption — every later step assumes the gates are real.



Rules for the commands themselves: `argv` is a **list**, never a shell string; call binaries by bare name so they resolve in the operator's environment; and put scope and strictness in your config files (`pyproject.toml`, `package.json`) rather than in the `argv`, so the gate and a local run can never disagree.

Part A — an existing codebase

Order: install + recon → wire the gates (rule zero) → set boundaries → interrogate → work.

A1. Recon before spending

```
cd /path/to/your/repo
uv run /path/to/sssf/.claude/skills/sssf/scripts/install.py
git add -A && git commit -m "Install SSSF"      # keep the tree clean, see A3
just scout "map this codebase: entry points, layers, test setup, build
           commands"
```

`just scout` is read-only. Its findings land in `adws/adw_data/sessions/<adw_id>/context_handoff/scout_findings.md`. Read them. You are checking whether the agent understood your repo before you let it change anything.

A2. Wire the gates (rule zero)

Point `test`, `lint` and `typecheck` at your real commands. Delete blocks you do not have — a `build` block running `echo` is a phantom check, and `run_quality()` reports it as passed. (That behaviour is in `adws/adw_modules/quality.py` — read `run_quality()` if you want to see it yourself rather than take this on trust.)

If your suite is slow, scope the test gate to a fast subset now and widen later. A gate nobody waits for gets disabled, and a disabled gate is a placeholder with extra steps.

A3. Set the boundaries

Two different mechanisms, and people conflate them:

- `tools` is a capability list — what the agent *can* do.
- `writes` is a boundary — what it *may change in the repo*, enforced after every call by diffing the tree.

Both are documented behaviour of the agent runner in `adws/adw_modules/` — grep for `writes` there to see the enforcement point.

For an existing codebase, tighten `writes` before your first build run. Give the documenter `docs/` and `**/*.md`. Give the planner `specs/`. Leave the builder unrestricted only if you are ready to review its whole diff.

`protected_files` is off-limits to everyone: keep `adws/adw_modules/`, `adws/adw_sssf_config/` and `adws/adw_*.py` in it, so no agent can edit the machinery that grades its own work. Add your CI config and any secrets path.

Sharp edge: the commit phase runs `git add -A` — it stages the *entire* working tree. Grep for `add -A` in `adws/adw_modules/` to confirm on your version. Start every run from a clean tree, or your unrelated work in progress lands in the agent’s commit.

A4. Interrogate before you plan


Same order as Part B, applied to code that already exists: interrogation comes before spec-writing, spec-writing before slicing.

Do this interactively in Claude Code, where your installed skills are available:

1. `/grill-with-docs` — interrogate your own intent against the scout findings from A1: what actually breaks today, what must not change, what you are assuming about the existing design. Writes the resolved vocabulary to `CONTEXT.md / docs/adr/` as you go.
2. `/to-spec` — synthesize that conversation into a spec. No interview; it already happened.
3. `/to-tickets` — slice the spec into per-phase tickets with `Blocked by:` edges.
4. Both (2) and (3) self-apply `ready-for-agent` by default — see Part D’s Filing “sharp edge” before committing their output.

Then hand one slice to SSSF. `just plan` is SSSF’s own synthesis step, and it is downstream of your thinking, not a replacement for it.

A5. The four jobs



Improve architecture. Do not hand this to a build workflow. Grill, spec and slice first (A4), then plan one slice, read the plan, then decide:

```
just plan "extract the payment provider calls behind one interface; no
          behaviour change"
# read specs/<adw_id>_*.md yourself – this is the cheap checkpoint
just sdhc "<same slice, once the plan is right>"
```

`just plan` runs the planner alone. Finding out the agent misunderstood your architecture costs one planner run instead of a full build plus a bad diff.

Find and fix bugs. The fix loop is a repair machine, so give it something red to repair. Reproduce the bug as a failing test first — by hand, or with a cheap agent — commit it red, then:

```
just sdhc "make the failing test in tests/test_x.py pass; do not modify the
          test"
```

The `do not modify the test` clause matters: the builder has write access, and deleting an assertion is the cheapest path to green. Constrain it explicitly rather than relying on virtue you have not asked for.

Implement a feature. The main line. Keep the slice narrow enough that one plan can express it:

```
just sdhc "add X; follow existing conventions; add tests covering A, B and C"
```

Cost scales with plan size. Measure your own first few runs with `just sessions` before you budget the rest.

Refactor into reusable abstractions. The most dangerous job, because “working” and “unchanged” are different claims. Refactoring is only safe behind tests that pin *current* behaviour, so:

1. Check coverage of the code you are about to move. If it is thin, spend a run adding characterisation tests **first**, and commit them.
2. Then refactor with an explicit no-behaviour-change constraint:

```
just sdlc "extract the discount rules in the cart-pricing module behind one
interface. No behaviour change: every existing test must pass
unchanged. Do not modify or delete any existing test."
```


The green suite is your evidence the refactor was safe. Without it you are not refactoring, you are rewriting and hoping.

A6. Read what happened

```
just sessions          # last runs, cost, status
just phases <adw_id>   # which phase did what
```

Check cost per run early and often, and build your own baseline from those numbers — no figure from someone else's repo transfers to yours. If one agent is eating most of your spend, try a cheaper model on it and compare both the price *and* the output it produces. Do not assume the expensive model plans better; run the same prompt through both on your own repo and read the two specs.

Part B — a new project



B1. Think before there is code (interactive, not SSSF)

SSSF is an execution engine. It is the wrong tool for deciding what to build.

Do this part **interactively** in Claude Code, where your installed skills are available. The order matters, and it is the reverse of what most people assume:

1. `/grill-with-docs` — interrogate relentlessly, round after round, until nothing important is left silently assumed. This happens **BEFORE** any spec exists — it is how you reach shared understanding. Writes the resolved vocabulary to `CONTEXT.md` / `docs/adr/` as you go, so it survives past this conversation.
2. `/to-spec` — synthesize that conversation into a structured spec. No interrogation here; it already happened.

3. `/to-tickets` — break the spec into per-phase tickets, each independently shippable and verifiable, with explicit `Blocked by:` edges.
4. Both (2) and (3) self-apply `ready-for-agent` by default — see Part D’s Filing “sharp edge” before committing their output.

Skipping the grill and starting at the spec produces a document that reads well and encodes assumptions nobody tested. The spec step is downstream of the thinking, not a substitute for it.

What the difference looks like. Same feature request — “let users export their data” — written two ways.

Straight to spec, no interrogation:

Add a data export feature. Users can export their data as CSV.
Acceptance: user clicks Export, receives a CSV file.

After a grill that asked which data, how much, who is allowed, and what happens when it is slow:

Export scope: only records the requesting user owns (confirmed: no team-shared records in v1).
Volume: 95th percentile account is ~40k rows → synchronous download times out.
Export is a background job; user gets an email link. Link expires in 24h.
Format: CSV only. JSON was requested and deferred — decided in the grill, not discovered mid-build.
Excluded: soft-deleted records, other users' PII in shared audit rows.
Acceptance: 40k-row account exports without timeout; a user cannot export a record they do not own; expired link returns 410.

The second is not longer for its own sake. Every extra line is a question that would otherwise have been answered by the agent guessing, mid-build.

Why still interactive: SSSF agents run with `--setting-sources ''`, so they cannot see the skills installed on your machine. Grep the agent launch args in `adws/adw_modules/` to confirm. That has always been true and remains true. The addition is the `skill_engineering:` config key: a *specific named* skill can be vendored into the repo and attached to one agent (see “Where the two layers sit”), so the two layers can meet —

but only for skills you choose on purpose, and only for the coding agents SSSF knows how to hand a system prompt to (currently all four it ships with). Your general interactive toolkit is still yours alone.

B2. Make it a repo with a walking skeleton

```
mkdir project && cd project && git init
```

Before installing SSSF, build the thinnest thing that runs and is tested — by hand or in one interactive session:

- one real module with one real function
- a test runner that works, with at least one passing test
- lint and typecheck configured

This is the target the factory needs. An empty repo gives the gates nothing to check, and you are back to placeholder-green.

```
git add -A && git commit -m "Walking skeleton"
```

B3. Install and wire (rule zero)

```
uv run /path/to/ssf/.claude/skills/ssf/scripts/install.py
# wire quality.py to your real commands, then prove each gate fails
just quality "gates wired"
git add -A && git commit -m "Install SSSF and wire the gates"
```

B4. Tune the roster

Start with everything on `claude_code` — no API key, uses your logged-in session. Then:

- `writes: []` on every read-only agent (scout, reviewer)
- documenter limited to docs paths
- planner limited to `specs/`
- cheap model on scout and documenter, stronger on planner and reviewer

- add your CI config to `protected_files`

B5. Run the slices, one at a time

Feed the slices from B1 in one at a time:

```
just sdlc "<one slice, stated as an outcome with its acceptance criteria>"
```

Review every diff before starting the next. The factory removes typing, not judgement. One slice per run keeps the blast radius small and the plan cheap.

The standing checklist

Before any run that writes:

☐

tree is clean (`git status`) — the commit phase stages everything

☐

gates wired, and each one watched failing at least once

☐

`writes` set for every agent that should not touch the whole repo

☐

`protected_files` covers CI config and the factory's own machinery

☐

you are on a branch you are willing to throw away

After:

☐

read the diff, not just the green banner

☐


`just sessions` — did it cost what you expected

☐

if a gate never went red across several runs, suspect the gate

Where the two layers sit

Skill vrending ships in the current version of SSSF: the vrending command and the `skill_engineering:` config key are both present. The layers meet on purpose, one skill at a time:




How it works. Each claim below is documented behaviour you can verify in `adws/adw_sssf_config/sssf.config.yaml` and the agent runner in `adws/adw_modules/` :

- **Vendor** the skill file into the repo with the vendoring command. The text now lives in your tree, versioned with your code.
- **Attach** it with the `skill_engineering:` config key on an agent in `sssf.config.yaml`. That agent's prompt carries the protocol.
- **Every coding agent, delivered differently.** `claude_code` and `pi` both take it as a real `--system-prompt` CLI flag; `agy` and `opencode` have no such flag, so they fold it into the user turn instead — weaker (advice inside the conversation, not a separate channel) but it arrives. Check `skill_engineering_applies()` in `adws/adw_modules/agents.py` to confirm which coding agents are covered on your version; a future coding agent added to SSSF needs a deliberate addition there before this key does anything for it.
- **Cost is visible.** Vendored skill text is real tokens on every call that agent makes, and the run reports it in the session record `just sessions` reads. Attach protocols you want, not every protocol you own.

Part C — going dark: bootstrap vocabulary once, then let the loop run headless

Parts A and B get you to `just sdlc` running safely. This part is what “dark factory” adds on top: minimal human intervention *across many features*, not just within one run. Two problems block that, and both have the same shape — something a human normally supplies mid-conversation has to be supplied *before* the headless loop starts instead, or the loop stalls or guesses.



Problem 1 — the interview a vendored planning skill wants doesn't have anyone to answer it

`wayfinder`'s "ask the user how to proceed" fallback is written for an interactive session — a human on the other end who answers back. Vendored under `skill_engineering:` and run by a headless node — any `coding_agent` — there is no one there. Best case the model role-plays both sides and you get a low-fidelity spec with none of the interview's real value; worst case it tries to prompt and the run hangs waiting for input that never comes.

Do not delete the interview — relocate it. The interview is genuinely how a project converges on shared vocabulary (what does "cancellation" mean here, is "account" the Customer or the User), and that can't be skipped just because a run is unattended. Split it into two phases that run in different modes:

- **Bootstrap, human-supervised, outside the headless loop.** When a feature introduces domain concepts your project hasn't named yet, run `/grill-with-docs` yourself — interactively, in whichever harness you're in. It composes `grilling` (the relentless interview) with `domain-modeling` (writes the resolved vocabulary into `CONTEXT.md` / `docs/adr/` — durable and citable, not prose trapped in one spec). Approve it, then **commit** `CONTEXT.md` / `docs/adr/` **before starting** `just sdlc` — scout and planner read them from the working tree, not from any live session. Skipping the commit makes the whole bootstrap step invisible to the pipeline.
- **AFK, headless, the common case.** If the request only composes concepts `CONTEXT.md` already names, skip bootstrap entirely — there's nothing left to interview about. `to-spec` already doesn't interview by design — "no interview, just synthesis," per its own frontmatter — so it runs unmodified (do not edit the vendored copy) and degrades to pure spec-formatting from already-agreed vocabulary. Add this to your `skill_engineering` -attached agent's own `system.md` (not the vendored skill file) so it knows what to do when the composed skill text below it says "ask the user":
 - `wayfinder` finds no fog → continue directly into `to-spec`, don't stop and ask how to proceed.
 - `to-spec` never prompts. If it hits a genuine ambiguity `CONTEXT.md` can't resolve, it says so in the plan and in `notes_for_next_agent`, proposes a best-guess term explicitly labeled provisional, and flags that a human should run the bootstrap interview before the plan is final — it does not guess silently and does not attempt to prompt.

- `to-tickets` 's “quiz the user” step never applies in-loop — the spec is always already in context from the prior step in the same composed prompt.

Two more overrides belong here, on top of the interview question — both `to-spec` and `to-tickets` self-apply the `ready-for-agent` triage label by default, which bypasses this pipeline's own `/triage` gate (Part D's “the judgment call stays interactive” depends on `/triage` being the only promoter):

- `to-spec` files at `Status: needs-triage`, not the `ready-for-agent` its own instructions say to apply. Skipping triage here would let a spec bypass this pipeline's own feasibility/compliance/redundancy judgment entirely.
- `to-tickets` files every ticket at `Status: needs-triage` too. Its own local-ticket-template uses bold `**Status:** / **Blocked by:**` lines — `adw_watch.py` 's `STATUS_RE / BLOCKED_BY_RE` tolerate up to two leading and trailing `* ( ^\{*{0,2} Status:\{*{0,2}... )`, added 2026-09-14 after four real portfinder tickets in bold sat invisible to the frontier scan, so bold-filed tickets are correctly matched today. What the regex still does **not** tolerate: backtick-wrapped ``Status:`` — that form is a real, still-open gap (see `docs/reference/the-queue/Frontier.md` and the queue lesson `06-going-dark / 07-the-queue` for the current state of this).

`grill-with-docs` should stay **out of** `skill_engineering/` — its `disable-model-invocation: true` is load-bearing, not an oversight. Vending it risks it landing in an agent's composed prompt and firing headless despite that flag, which is exactly the failure this split exists to prevent.

Problem 2 — nothing stops a plan from re-implementing what already exists

Neither `prompt_engineering/` nor any shipped `skill_engineering/` skill tells scout or the planner to check what's already built before scoping new work. Left unfixed, every feature costs a full plan-and-build cycle even when 80% of it already exists — the opposite of minimal diff.

Add two lookups to **scout's** `system.md`, before any other search:

1. **Glossary first**, if `CONTEXT.md / docs/adr/` exist (see Problem 1) — note any request term missing or conflicting; that's a signal for the planner, not something scout resolves itself.

2. **A structured knowledge source second**, if your project maintains one — an OKF-style knowledge bundle, a `CONTEXT.md` -linked concept map, or similar. Report what's *already implemented* and how, cited by concept path, not line numbers, which drift.
3. **A structural code search third** (e.g. `ast-grep`) to confirm the knowledge source's claims against actual code. Documentation can lag; ground truth doesn't. When they disagree, trust the code and flag the doc as stale.

Then in **planner's** `system.md` : read scout's findings first, and scope the plan to the delta only — anything scout already marked implemented is explicitly out of scope, not silently re-touched.

None of this needs a new SSSF agent role or a new `skill_engineering` entry — `ast-grep` and a knowledge-source lookup are tool-usage patterns baked into scout's own instructions, the same way `writes` and `protected_files` already are. They don't belong in the same composition list as `wayfinder` / `to-spec` / `to-tickets` / `tdd` , which are planning methodologies.

The grill → triage handoff is manual, on purpose (accepted, not a gap to close)

`/grill-with-docs` ends at a committed `CONTEXT.md / docs/adr/` update; it does not hand `/triage` anything, and neither does upstream Pocock's own flow (see `.okf/pocock-skills/naming-drift.md` and the integration plan's finding F2 — checked against the current upstream skill set too, not just the version this playbook was originally written against; nothing there bridges it either). The bootstrap/AFK audit's finding 3.1.3 flagged this as the one place in the whole pipeline where forgetting a step loses work silently: a grilled-and-settled feature that never reaches `/triage` never enters the queue, and nothing lists “settled but unfiled” work.

Decision: accept this as a manual step, not a mechanism to build. Closing it would mean a thin wrapper skill (run `/grill-with-docs` , then `/to-spec` , then remind `/triage`) — evaluated and deliberately not built, since it adds a skill to maintain for one reminder a checklist line covers just as well. The checklist item below is the actual guard.

Definition of done, extended

Everything in the standing checklist above, plus:



If bootstrap ran, `CONTEXT.md / docs/adr/` are committed and predate the `just sdlc` run that used them.



The plan explicitly excludes anything scout marked as already implemented — check `specs/<adw_id>*.md` names what it's *not* doing, not just what it is.



Your knowledge source (OKF, concept map, whatever you use) got updated by the documenter stage if the feature changed anything it describes — otherwise the next feature's discovery step is reading stale claims.



If this bootstrap session ran because of a fresh grilling (`/grill-with-docs` or `/improve-codebase-architecture`), its outcome is filed — run `/triage` before leaving the session. Nothing bridges this automatically (see above); a settled-but-unfiled grill is invisible to the queue, not just slow to reach it.

Part D — the queue: from `ready-for-agent` to shipped

Part C gets one feature through the headless loop unattended. This part is what turns that into a standing queue: an engineer files or triages work whenever they want, and something — `just watch`, watching — picks it up and ships it without anyone manually running `just sdlc` per item.


A note on paths in this part. `docs/agents/issue-tracker.md`, `.scratch/`, and `sssf.config.yaml` below all mean paths **in the repo you are adopting SSSF into**, once `install.py` has stamped it and you've set up the tracker — not paths in this skill-source fork, which has none of them.

The whole chain, named once

1. `/grill-with-docs` — interactively, commit `CONTEXT.md / docs/adr/`.
2. `/to-spec` — synthesize the spec.
3. `/to-tickets` — slice into tickets.
4. `/triage` — the only thing allowed to promote a ticket to `ready-for-agent`. (Steps 2–4: see Filing's “sharp edge” — running (2)/(3) interactively self-promotes past this gate by default; fix that there, not here.)

5. Commit; run the Mandatory checkpoints.

6. `just watch` .



Filing: less manual than it looks, but one real gap remains

Correction (v4.5): earlier versions of this section claimed nothing writes `issues/NN-slug.md` . Traced against a real repo's filed tickets and that's wrong for that file specifically — `/triage` 's own "apply the outcome" step, for `ready-for-agent` , is "post an agent brief comment," and

`docs/agents/issue-tracker.md` says “when a skill says ‘publish to the issue tracker’ → create a new file.” On the local-markdown tracker, `/triage` posting its agent brief on a not-yet-tracked item **is** the file-creation event — verified by content only `/triage`’s documented flow produces (a `.out-of-scope/` prior-rejection check, a redundancy check, the exact category/state role vocabulary) showing up inside filed issue files that were never touched by any other skill. So: `/triage` **files** `issues/NN-slug.md` **itself**, one full pass per item, as long as it’s handed something to triage.

The remaining gap is narrower and upstream of that: `grill-with-docs` (`grilling` + `domain-modeling`) and `improve-codebase-architecture`’s grilling loop both end at updated `CONTEXT.md` / `docs/adr/` — **neither hands** `/triage` **anything**. Something still has to turn “we settled on this” into the first description `/triage` can act on. That’s the actual manual step: you (or an assisting agent, same interactive session) either invoke `/triage` directly with the settled description in natural language (no file needed first — triage creates it), or write a minimal stub yourself if you want the file to exist before triage runs.

Pick the shape, per `docs/agents/issue-tracker.md`’s own convention:

- **New feature** → `.scratch/<feature-slug>/spec.md` — a spec (problem, solution, user stories). `/to-spec` writes this directly; no manual step here at all once you’ve run it.
- **Addition to an existing feature** → `.scratch/<existing-feature-slug>/issues/NN-<slug>.md`, `NN` the next free number in that feature’s `issues/` dir. Numbers are scoped per feature, not global — a `Blocked by: 01` in one feature never refers to another feature’s `01`. Hand `/triage` the description; it files the ticket.

Sharp edge running either of these interactively: they self-promote past `/triage`. `/to-spec` and `/to-tickets` both apply the `ready-for-agent` triage label themselves by default (their own “no need for additional triage” / “the tickets are agent-grabbable by construction” instructions) — this pipeline’s headless planner overrides that (see Part C’s AFK bullet list), but running either skill **by hand** gets you their unmodified behavior: an **untriaged** `ready-for-agent` ticket (bold or plain — both forms are visible to `adw_watch.py`’s frontier scan) — `just watch` will build it without anyone having judged feasibility, redundancy, or compliance first, exactly the check Part D’s “judgment call stays interactive” section depends on `/triage` providing. After running either skill interactively, either hand its output to `/triage` before committing, or hand-edit the `Status:` line to `needs-triage` and confirm the file uses plain (not bold) `Status:` / `Blocked by:` lines — same shape the Mandatory checkpoints below already check for.

If you write a stub yourself (rather than handing `/triage` a bare description), it must contain what `/triage` and `just watch` both actually parse, not just prose for a human:

```
# Short title
```

```
Status: needs-triage
```

```
Whatever the grilling settled on — the decision, the reasoning, any
CONTEXT.md terms or ADRs it touches.
```

- `# Title` — a heading.
- `Status: needs-triage` — exact shape (`Status:` + one of the five canonical triage states). `adw_watch.py`’s frontier logic and `/triage`’s state machine both key off this line; get the string wrong and both silently ignore the file.
- `Blocked by: NN, NN` — optional, only if it depends on another ticket already in the same feature’s `issues/` dir.

Then commit it. `/triage` and `just watch` read from the working tree, not from your conversation — an uncommitted file is invisible to a fresh session, same reasoning as committing `CONTEXT.md` /ADRs before the headless loop starts.

Two owners of one file — protect the queue’s extensions on re-sync

`docs/agents/issue-tracker.md` itself has two owners. `/setup-matt-pocock-skills` (not part of SSSF) writes its base from its own canonical local-tracker template. `just watch` then depends on that same file also carrying SSSF’s own additions — most concretely, the `claimed / resolved` states `adw_watch.py` writes, which the base template’s own wayfinding-operations prose doesn’t name. Nothing currently merges these automatically: if `/setup-matt-pocock-skills` is ever re-run to refresh the base, it has no way to know an SSSF extension was layered on top, and a whole-file regeneration silently drops it. Wrap anything SSSF added or reworded in `<!-- sssf:queue-extension --> ... <!-- /sssf:queue-extension -->` markers — same idiom `vendor_skill.py` uses for its provenance headers — so a re-sync has something greppable to diff against rather than relying on someone remembering. This is a diff aid, not automatic reinjection; there is no tooling yet that restores a dropped block on its own.

Mandatory checkpoints — verify the filing actually happened

Every one of these has failed silently on a real repo this playbook was built against: a skill claims to have written something, and the file isn’t where the claim said, or isn’t committed, or the status string doesn’t match what the state machine expects. Run this after `/to-spec`, `/to-tickets`, or `/triage` — before assuming the queue can see the work:

```

# 1. The feature directory actually exists
ls .scratch/<feature-slug>/                                # expect: spec.md, plan.md (if
                                                            # planned), issues/ (if any
                                                            # scope was deferred)

# 2. spec.md / plan.md exist if the skills that claim to write them ran
test -f .scratch/<feature-slug>/spec.md && echo "spec.md: present"
test -f .scratch/<feature-slug>/plan.md && echo "plan.md: present"

# 3. Every issues/ file has a parseable Status: line, exact shape
grep -L "^Status: " .scratch/<feature-slug>/issues/*.md      # empty output =
good;                                                         # any listed file
                                                             # invisible to
is                                                            # /triage and
                                                             # adw_watch.py's
both                                                         # frontier scan

# 4. Status values are one of the five canonical roles – a typo here is
#   silent, not an error
grep -h "^Status: " .scratch/<feature-slug>/issues/*.md \
| sort -u             # every line must be one of: needs-triage, needs-info,
                      # ready-for-agent, ready-for-human, wontfix

# 5. Nothing is sitting uncommitted – invisible to a fresh session, same
#   failure mode as an unwired gate reporting a phantom PASS
git status --short .scratch/<feature-slug>/                  # expect: empty

# 6. Blocked-by references actually resolve within the same feature
grep -h "^Blocked by: " .scratch/<feature-slug>/issues/*.md
# then confirm each NN referenced has a matching <feature-slug>/issues/NN-*.md

```

If any of these come back wrong, don't proceed to `just watch` — a `Status:` typo or an uncommitted file doesn't error, it just makes `adw_watch.py`'s frontier scan silently skip the item forever, the same class of manufactures-confidence failure rule zero exists to catch at the gate level.

Check 6's sharp edge, found live (since fixed): an earlier `adw_watch.py` captured everything after `Blocked by:` and split it on commas, treating every fragment as a blocking ticket number. A hand-written value like `none (API exists, see spec.md)` — a comma *inside* the parenthetical — split into two fragments, neither matching a real sibling ticket, so `is_unblocked()` returned `False` and the ticket sat at `ready-for-agent`, invisible to the frontier scan, with `just watch --once` reporting "queue empty." The parser now strips any parenthetical from the whole value *before* the comma split, and drops bare `none`, so that

exact line is read as “no blockers” — covered by a regression test. The simpler habit still holds: if a ticket has no real blocker, omit the `Blocked by:` line entirely; a value with nothing to parse is the one that can never mis-parse.

The judgment call stays interactive, on purpose

It’s tempting to have `just watch` itself analyze a raw issue’s feasibility, compatibility with what’s already live, and compliance/security against `CLAUDE.md / AGENTS.md`, `CONTEXT.md`, your knowledge source, and `docs/` /plan files — then decide whether to build it. Don’t: that recreates the exact failure Part C already fixed once: `wayfinder`’s “ask the user how to proceed” fallback had no one to answer it headless (Problem 1). An unattended loop making its own security/compliance judgment calls has no one to catch it when it’s wrong, either.

The `triage` skill (if installed) already does this analysis — feasibility, redundancy against existing implementation, `.out-of-scope/` prior-rejection checks, compliance — but it is explicitly interactive: it waits for maintainer direction at every step and never runs unattended. Its terminal state for approved work, `ready-for-agent`, means a durable **agent brief** has been posted to the issue. That’s the same pattern as `grill-with-docs`: the judgment happens once, by a human, and only the *result* — not the judgment process — becomes something a headless loop can safely consume later.

Current `triage` (`disable-model-invocation: true`) also covers **external pull requests** as a request surface — “a PR is an issue with attached code,” same roles, same states, same machine — if the tracker config marks them in scope. Off by default, so nothing here changes until you turn it on. Note the ceiling this hits today, though: a PR only exists on a GitHub/GitLab tracker, and `adw_watch.py` currently refuses those outright (exit code 2, local-markdown only — see its own module docstring). Triage can flip a PR to `ready-for-agent`; `just watch` has nothing that would ever pick it up until the GitHub/GitLab tracker support the watcher’s docstring already flags as unbuilt actually lands.

So: `just watch` **only ever picks up issues already in `ready-for-agent` state**. Everything that decides whether something is safe, compatible, and compliant to build happens upstream of that label, interactively. Scout and planner’s Part C wiring (glossary/OKF/ast-grep discovery, delta-scoping) still runs inside every `just sd1c` call underneath this — defense-in-depth, not a substitute for the triage gate.

The queue mechanism already exists — reuse it

Don't build a new work-item format. `docs/agents/issue-tracker.md` 's "Wayfinding operations" section (written by `/setup-matt-pocock-skills`, see Part A/B) already defines everything a queue needs: a **map** + numbered **child** files, a `Blocked by:` line, a **frontier** rule ("scan for files that are open, unblocked, and unclaimed; first by number wins"), and **claim/resolve** semantics. That's dependency-ordered work-item tracking, full stop — built for a human working research tickets, but the mechanism doesn't care who's claiming.

`just watch` reuses it verbatim: same frontier scan, same claim-before-work, same resolve-after-work, just performed by an agent instead of a human.

This also depends on a `to-tickets` slice's individual **phases** (Part A5/B1) each having their own atomic, independently-pickable file — not just the feature's `spec.md`. `to-tickets` natively files one ticket per phase, with a `Blocked by:` line expressing phase ordering (Phase 2 blocked by Phase 1) the same way a wayfinder ticket blocks on another, so a queue can claim work safely one phase at a time.

What `just watch` does, concretely

A polling loop, tracker-aware (reads `docs/agents/issue-tracker.md` to know whether issues live on GitHub or under `.scratch/`):

1. **Scan** for `ready-for-agent` issues — `gh issue list --label ready-for-agent` or `grep Status: ready-for-agent` across `.scratch/*/issues/*.md`.
2. **Frontier**: filter to unblocked (every `Blocked by:` target already resolved) and unclaimed; oldest/lowest-numbered first.
3. **Claim**: set `Status: claimed` before touching anything else and commit that write immediately — same as wayfinder. This is ordering, not a lock: a second `just watch` whose scan runs after the first's claim commit sees the ticket taken, but two runs scanning at the same instant, before either commits, is a gap the design does not close. Run one watcher per repo.
4. **Dispatch**: extract the agent brief, run the full SDLC chain on it — `adw_simple_sdlc.py` (`planner -> builder -> reviewer -> revision loop -> documenter -> commit`), not the lighter `plan_build_test` chain `just sdlc` wraps. An unattended dispatch has nothing else checking it besides the test gate passing, so it gets the review/revision loop a supervised manual run could otherwise skip.
5. **Resolve**: on success, close the issue and append a pointer (session id, commit, cost) — mirroring wayfinder's "append a context pointer to the map." On failure, flip to

`ready-for-human` with a comment explaining what broke. Never retry silently — a loop that keeps re-attempting the same failure burns money without anyone finding out until later.

6. **Loop** on an interval, or — more composable — run once-and-exit, invoked periodically by `cron / launchd / CI` outside SSSF's own scope, rather than the factory growing its own daemon.

Definition of done, extended again

Everything in Part C's definition of done, plus:



Every issue `just watch` picked up was in `ready-for-agent` state when it claimed it — never a raw or `needs-triage` issue.



A failed run is `ready-for-human` with a comment explaining what broke, not silently retried or left `claimed` forever.



Each `to-tickets` phase intended for the queue has its own issue file with `Blocked by:` expressing phase order, not just a shared `plan.md`.

What this playbook does not claim

- That agent output needs no review. Agents fabricate confidently, including about their own work. Read the diff before you keep it.
- That vendoring a skill makes an agent follow it. Attaching `tdd` puts the protocol in the prompt. Nothing verifies the agent worked test-first — the gates judge the outcome, not the process. That distinction is the point.
- That any cost figure transfers to your repo. Language, suite size, and repo shape dominate. Measure your own on a throwaway branch before budgeting.
- **Resolved 2026-09-13** (was: not yet run live). The full chain — `grill-with-docs` → `to-spec` / `to-tickets` (inline, in this run — see the v4.15 changelog row) → `triage` → Mandatory checkpoints → `just watch` 's current dispatch target (`adw_simple_sdlic.py` : planner → builder → reviewer → revision loop → documenter → commit) — has now been run live, once, end to end (a downstream project, device-management-ui, `adw_id: c1eff7ac`): 10/10 phases passed, reviewer approved 12/12 plan requirements and 6/6 acceptance criteria, real cost (\$4.42, 4.3M tokens), real

commits (plan, build, docs, each its own commit), ticket auto-flipped to `resolved`. One real defect surfaced along the way and was fixed before dispatch succeeded (the free-text `Blocked by: gap, v4.15`) — the run did not go perfectly on the first try, and this playbook does not claim it will for you either. One live run is evidence the chain *can* work end to end, not that it reliably will across other repos, request shapes, or failure modes this single pass didn't happen to hit.

Cheat sheet — the five vendored skills

Quick reference for the standard roster (`wayfinder`, `to-spec`, `to-tickets`, `tdd` on planner; `code-review` on reviewer). Full detail is in “Where the two layers sit” and Parts C/D above; this is the lookup table.

Which role uses which skill, and how:

Skill	Vendored to role	Interactive or headless?	Output → handoff location
<code>wayfinder</code>	planner	Headless in-pipeline (no-fog fallthrough forced — see Part C, Problem 1)	decides fog/no-fog; no separate artifact
<code>to-spec</code>	planner	Headless (never interviews, by design)	spec → specs/ <adw_id>_<slug>.md
<code>to-tickets</code>	planner	Headless (quiz skipped; forced <code>needs-triage</code> , plain <code>Status: / Blocked by: lines</code>)	tickets → <code>.scratch/</code> <feature>/issues/ NN-slug.md (only on a fresh, non-dispatched request — see the planner's dispatch-detection guard)
<code>tdd</code>	planner	Headless (pure methodology)	shapes <code>plan.md</code> 's phases red→green

Skill	Vendored to role	Interactive or headless?	Output → handoff location
<code>code-review</code>	reviewer	Headless (the planner's <code>writes: specs/ grant</code> means <code>plan.md</code> always lands in <code>context_handoff_dir</code> , which the reviewer's own prompt reads before falling back to <code>prompt</code> — see the trap below)	verdict → <code>context_handoff/review.md</code>

Trap: `code-review` has two interactive fallbacks this table does not fully close.

Vendoring the skill composes its text onto the reviewer's existing `system.md / user.md` (SSSF's own built-in review methodology); it doesn't rewrite the skill's own "Process." That real process says "pin the fixed point... if they didn't specify one, ask for it" and "if nothing is found, ask the user where the spec is" — both written for an interactive session. SSSF's reviewer prompt already supplies an equivalent (reads `plan.md` from `context_handoff_dir`, falls back to `prompt`), which happens to satisfy the spec-source question in practice, but nothing documents that this is deliberate coverage for `code-review`'s own ask-conditions, and the fixed-point question is not addressed at all. Unlike `wayfinder`'s fallback (Part C, Problem 1, treated in full), this one has never been traced end to end — treat it as unverified, not resolved.

Not vendored — interactive-only, run by a human outside any agent:

Skill	Who runs it	Mode	Output
<code>grill-with-docs</code>	you (Stage 0, only for new vocabulary)	Interactive	<code>CONTEXT.md / docs/adr/</code>
<code>triage</code>	you (Part D's judgment gate)	Interactive	flips a ticket's <code>Status: to ready-for-agent</code> (or <code>wontfix / ready-for-human</code>)

Handoff format and location, stage → stage:

From → To	Format	Location
you → planner	raw request text	<code>just sdhc "<text>" arg</code> , or a claimed <code>issues/NN-slug.md</code> 's body
scout → planner	findings doc	<code>context_handoff/scout_findings.md</code> (runtime, gitignored)
planner → builder	plan + spec	<code>context_handoff/plan.md</code> (runtime) + <code>specs/<adw_id>*.md</code> (committed)
builder → reviewer	code diff	working-tree diff, no separate file
reviewer → documenter/ builder	verdict	<code>context_handoff/review.md</code> (runtime)
documenter → done	doc	<code>app_docs/</code> , <code>docs/</code> (committed)
watcher → planner (queue only)	filed ticket, verbatim	<code>.scratch/<feature>/issues/NN-slug.md</code>

Everything under `context_handoff/` is runtime state (gitignored, one `adw_id`'s own working set); everything else in the right column is what survives to the next run and what a human reads afterward.

Cheat sheet — the extended Pocock+ToB roster (validated downstream)

The five-skill cheat sheet above is SSSF's own shipped default. A downstream adopter (`weather-report`) extended it with a cross-ecosystem research pass — every role's real `system.md` / `user.md` checked against both the full Pocock catalog and the full Trail of Bits catalog (`trailofbits/skills` , ~44 plugins/88 `SKILL.md` files), using the six-check compatibility framework in `docs/reference/agent-configuration/Skill-compatibility-checklist.md` . This is downstream-repo config, not a change to SSSF itself — SSSF stays

skill-agnostic (see “Where the two layers sit”); recorded here because the pattern and its two validated runs are reusable evidence for the next adopter who wants to extend their own roster past the five defaults.

The roster, in composition order (current, 2026-10-07):

Role	<code>skill_engineering</code> list	New vs. the five-skill default
planner	<code>codebase-design-planner</code> , <code>tdd-planner</code> , <code>sharp-edges-planner</code>	Replaces <code>wayfinder</code> / <code>to-spec</code> / <code>to-tickets</code> entirely (see “Re-armed” below) — all three files are planner-specific trims, not the shared copies other roles use
builder	<code>codebase-design</code> , <code>tdd</code> , <code>property-based-testing</code>	+ <code>codebase-design</code> , <code>property-based-testing</code> (ToB) — builder has no default skill wiring in the five-skill table at all
scout	<code>codebase-design</code>	+ <code>codebase-design</code> (reference only)
reviewer	<code>sharp-edges</code> , <code>differential-review</code> , <code>code-review</code>	+ <code>sharp-edges</code> , <code>differential-review</code> (both ToB), ahead of the existing <code>code-review</code>
documenter	<code>writing-for-agents</code> , <code>pr</code>	+ both — documenter has no default skill wiring in the five-skill table either

The planner roster was re-armed once already, after the first version of this roster shipped with `codebase-design`, `tdd`, `wayfinder`, `to-spec`, `to-tickets`, `sharp-edges`. A 3-round opus-vs-fable debate (the adopting repo’s own `docs/agents/planner-roster-debate.md`) found `wayfinder` / `to-spec` / `to-tickets` were dead text on every queued ticket: the human bootstrap chain (`grill-with-docs` / `wayfinder` → `to-spec` → `to-tickets` → `/triage`, run interactively, outside this pipeline) already does everything those three skills’ text describes, before a ticket ever reaches `ready-for-agent` — the planner’s own dispatch-detection (see Part C, Problem 1) was already skipping them on every real queue

dispatch, so composing them was pure overhead with a live downside: `opencode`'s own skill auto-discovery (see the harness note below) could still resolve `wayfinder`'s own "ask" text from disk even after removal from this list, if the harness-level fix hadn't also shipped.

`codebase-design` is reference-only on every role it touches (shared seam vocabulary, no interactive asks, nothing to override). The other two skills on the planner, plus `property-based-testing` on the builder, each needed a real `system.md` amendment, because each has one interactive fallback the compatibility framework would flag:

- **planner + `tdd`** : `tdd`'s "write down the seams under test and confirm them with the user" never applies headless. Added a bullet — name the seams from the request/`scout_findings.md` yourself, record them as a "Seams" list in `plan.md`. Also struck `tdd`'s own "call the Skill tool with `codebase-design`" line — composition order already puts `codebase-design` first, and under the harness note below that call could otherwise resolve a stale global copy instead of this repo's vendored one.
- **planner + `sharp-edges`** : its own Phase 4 ("write minimal code demonstrating the footgun") assumes a code-writing role; the planner only writes `plan.md`. Amended to record each sharp edge as a named test the builder must write at a seam in `plan.md`'s own Seams list, not as a standalone report.
- **builder + `property-based-testing` (ToB)**: the skill says to "let the user decide" before adding a property-testing library as a new dependency. Added a headless override — reuse an existing `proptest` / `quickcheck` / `hypothesis` -class dependency if one is already named; otherwise add the smallest standard choice as a dev-only dependency and say so plainly in `notes_for_next_agent`, never ask. A companion scope-discipline line was added separately: don't add an inverse/decoder/round-trip test the ticket itself didn't ask for — a query/filter-shaped ticket has no meaningful inverse, and reaching for one anyway is headless scope creep dressed up as thoroughness.
- **reviewer + `sharp-edges` / `differential-review`** (both ToB): the reviewer's own "not your job: generic style opinions" line would put both skills' real findings out of scope by accident. Amended it to carve out footguns/security regressions as in-scope evidence (not style opinions), reported the same way as spec requirements, blocking only when a reasonable engineer would actually block a merge. Also redirected `differential-review`'s own "always generate a report file" instruction at the existing `review.md`, so the reviewer writes one report, not two.

A harness-level leak, found by the same debate, applies beneath all of this: the `opencode` coding-agent interface auto-discovers skills from both `~/.claude/skills/` (global) and the project's own `.claude/skills/` — independent of this harness's `skill_engineering:` composition — and injects the operator's own `~/.claude/CLAUDE.md` into every prompt. This

means removing a skill from a roster entry's `skill_engineering:` list does **not** make it unreachable: the skill tool can still resolve it from disk if a role ever calls it, and an operator's unrelated global steering file (persona rules, unrelated project conventions) rides into every prompt uninvited. The fix lives in the `opcode` coding-agent module itself (`OPENCODE_DISABLE_CLAUDE_CODE_SKILLS=1`, `OPENCODE_DISABLE_CLAUDE_CODE_PROMPT=1` on the subprocess env, scoped to that module only), not in any `sssf.config.yaml` — but it is load-bearing for every claim this section makes about what a roster's composed `skill_engineering:` list actually controls, for any role running on `coding_agent: opcode`. Confirmed live (not just inferred): a direct `opcode run --print-logs` call showed 33 "duplicate skill name" warnings pulling in `grilling / domain-modeling / to-tickets / triage` from disk and 72 tools/skills registered without the two env vars; zero skill-path log lines and 32 registered with them.

ToB skills need one more step this table's Pocock skills don't: `compose()` ships exactly one vendored file, but a ToB `SKILL.md` often splits its real substance across `references/*.md` siblings. `sharp-edges`, `differential-review`, and `property-based-testing` were each hand-flattened (the `SKILL.md` plus only its relevant reference files — e.g. `sharp-edges` dropped 11 other-language reference files, keeping only `lang-rust.md` and the language-agnostic ones) into one composite file, carrying a CC-BY-SA-4.0 attribution comment block naming the exact upstream commit, before being vendored through the normal `vendor_skill.py --as <name>` path. This is check 7 of the compatibility framework (sibling-file dependency — not yet folded into the checklist page itself as of this writing).

Validated in four independent live runs — the first two via direct `just simple-sdlc "<text>"` CLI dispatch, the last two via the real `just watch` queue path (see below for why that distinction matters). **Correction on run 1:** only run 2 used the promoted production `sssf.config.yaml` (commit `3c7873f`); run 1 predates that promotion and ran under a separate `sssf.config.experiment-pocock-tob.yaml` — its clean result is what led to the promotion, not evidence gathered *after* it.

#	adw_id	Dispatch	Ticket	Phases	Tests	Cost
1	a1fba942	CLI, pre-promotion config	PBT for a coordinate parser (<code>parse_coords</code>)	10/10	38/38 (4 new)	\$1.16
2	aca5e802	CLI, production config	Country filter on <code>GET /ports</code>	10/10	109/109 (5 new)	\$0.78

#	adw_id	Dispatch	Ticket	Phases	Tests	Cost
3	2480c78a	just watch , re-armed planner	Bug: NMEA minute carry (format_nmea_lat / lon)	10/10	7/7 nmea, 22/22 route_	\$1.15
4	380a46ec	just watch , re-armed planner	Bug: NaN edge weight in from_custom	10/10	27/27	\$0.51

None of the four hit a revise loop. Runs 1–2: reviewer produced one `review.md` with a distinct “Sharp-Edges Analysis” section separate from the spec-conformance checklist (the style-opinions amendment holding — new findings got a real reporting slot, without turning every observation into a blocker); builder added its new test-only dependency (`proptest`) with zero interactive pause. `property-based-testing` is expected to help mainly on codec/parser/numeric code and add little elsewhere; run 2 (a straightforward query-filter addition, not parser-shaped) still passed clean, but with less surface for that specific skill to add value on.

Runs 1–2 did not go through just watch ’s queue — a real, separate gap from the roster question. A queued ticket’s body already carries a `Status:` line, which the planner’s dispatch-detection treats as “already filed,” skipping `wayfinder / to-spec / to-tickets` entirely on the *original* roster (they’re gone from the *re-armed* one regardless). A real consequence surfaced: dispatching by direct CLI never flips the ticket’s own `Status:` line to `resolved` (only `just watch` ’s claim/resolve cycle does that), so a ticket built this way sits at `ready-for-agent` and gets rebuilt by the next real `just watch` claim unless someone updates its status by hand — this happened for real on run 2’s own ticket, caught and fixed by hand afterward. **Runs 3–4 close this gap:** both are real `just watch --once` dispatches, against the re-armed planner roster, each claiming, building, reviewing, documenting, and resolving a filed bug ticket with no manual status intervention needed — the watcher’s own commits (`claim` → `resolve`) did the bookkeeping correctly both times. Full trace of the gap and its fix: `weather-report/ docs/agents/bootstrap-afk-pocock-tob-research.md` and `docs/agents/planner-roster-debate.md` .

Version history

Version	Date	Changes
4.22	2026-10-07	Updated the Pocock+ToB cheat sheet for the downstream adopter's planner re-arm (3-round opus-vs-fable debate, <code>docs/agents/planner-roster-debate.md</code>): the planner's roster table was still showing the original <code>codebase-design</code>, <code>tdd</code>, <code>wayfinder</code>, <code>to-spec</code>, <code>to-tickets</code>, <code>sharp-edges</code> list, stale since the adopter replaced <code>wayfinder</code> / <code>to-spec</code> / <code>to-tickets</code> entirely with planner-specific trims (<code>codebase-design-planner</code> , <code>tdd-planner</code> , <code>sharp-edges-planner</code>) after finding them dead text on every queued ticket. Added a harness-level leak this same debate found and live-confirmed: <code>opencode</code> 's own skill auto-discovery and <code>CLAUDE.md</code> auto-injection bypass this harness's <code>skill_engineering: composition</code> entirely unless two env vars are set in the coding-agent module itself — load-bearing for every claim this section makes about what a roster controls. Added two more validated live runs (3, 4) that close the “neither run went through <code>just watch</code> 's queue” gap v4.21 flagged — both are real <code>just watch --once</code> dispatches against the re-armed roster, resolving two filed bug tickets with correct automatic status bookkeeping. Narrowed the results table from 8 to 7 columns after the PDF regeneration's own visual QA caught the Cost column clipping off the page edge in the rendered PDF — the stale content would have gone unnoticed without that regeneration.

Version	Date	Changes
4.21	2026-10-07	Corrected v4.20's own claim within hours: run <code>a1fba942</code> was cited as evidence gathered against the promoted production <code>sssf.config.yaml</code> , but it actually ran under the pre-promotion experiment config <code>(sssf.config.experiment-pocock-tob.yaml)</code> — only run <code>aca5e802</code> ran post-promotion. Found by a follow-up opus research pass cross-checking commit timestamps. Also added a caveat the same research surfaced: neither run went through <code>just watch</code> 's queue, so the queue path's <code>wayfinder / to-spec / to-tickets</code> dispatch-detection skip (which fires on any ticket body carrying a <code>Status:</code> line, old or new roster alike) was never actually exercised by either validated run — and dispatching by direct CLI instead of the queue leaves the ticket's own <code>Status:</code> line stuck at <code>ready-for-agent</code> , risking a redundant rebuild on the next real <code>just watch</code> claim. Full trace: <code>weather-report/docs/agents/bootstrap-afk-pocock-tob-research.md</code> .

Version	Date	Changes
4.20	2026-10-07	Added a new “Cheat sheet — the extended Pocock+ToB roster” section documenting a downstream adopter’s (<code>weather-report</code>) cross-ecosystem skill-engineering extension: codebase-design/tdd/sharp-edges on planner, codebase-design/tdd/property-based-testing on builder, sharp-edges/differential-review/code-review on reviewer, codebase-design on scout, writing-for-agents/pr on documenter. Documents the three required <code>system.md</code> amendments (planner’s tdd seam-naming override, builder’s headless PBT-dependency override, reviewer’s sharp-edges/differential-review scope amendment), the ToB sibling-file-flattening step (compatibility check 7, not yet folded into the checklist page), and two independent live <code>just simple-sdlc</code> runs against the roster as the adopting repo’s production default (not an experiment config) — both 10/10 phases, no revise loop. This is downstream-repo config, recorded here as reusable evidence, not a change to SSSF’s own skill-agnostic shipped default.
4.19	2026-10-06	Fixed a wrong citation in the Cheat sheet: <code>code-review</code>’s headless-safety was attributed to the reviewer’s <code>writes: specs/</code> grant, but that permission belongs to the planner, not the reviewer (whose own config is <code>writes: []</code>, read-only). The real mechanism is the planner’s write landing <code>plan.md</code> in <code>context_handoff_dir</code>, which the reviewer’s prompt reads. Found by checking a freshly-pulled copy of the Pocock skills against this playbook’s claims. Also flagged, not yet resolved: <code>code-review</code>’s own two interactive ask-fallbacks (fixed point, spec source) have never been traced end to end for headless safety the way <code>wayfinder</code>’s was.

Version	Date	Changes
4.18	2026-09-25	Two corrections surfaced by a teacher/student assessment round against the docs. (1) The Filing section's <code>Blocked by:</code> comma-split sharp edge is now described as fixed — <code>adw_watch.py</code> strips parentheticals before splitting and has a regression test — matching Chapter 7's filing-checkpoints lesson, which already said so. (2) Part D's "Claim" step no longer says concurrent watchers "never double-pick": claim-then-commit is ordering, not a lock, and a same-instant scan race is an acknowledged gap — matching Chapter 7's <code>just watch</code> lesson and <code>adw_watch.py</code> itself.
4.17	2026-09-14	Added a "Cheat sheet" section: three lookup tables for the five vendored skills (role/mode/output per skill), the two interactive-only skills (<code>grill-with-docs</code> , <code>triage</code>), and the handoff format/location between every conveyor stage (scout→planner→builder→reviewer→documenter, plus the queue's watcher→planner path). Also fixed the frontmatter <code>version:</code> field, which had drifted two versions behind the version-history table's own top row (said 4.14 while the table already listed 4.16) — bumping this same edit closes that gap rather than leaving it to compound further.
4.16	2026-09-13	B1 closed out: <code>just watch --once</code> succeeded live end to end after the v4.15 fix (a downstream project, device-management-ui, <code>adw_id: c1eff7ac</code> — 10/10 phases, reviewer approved 12/12 plan requirements and 6/6 acceptance criteria, real commits for plan/build/docs, ticket auto-resolved). Updated "What this playbook does not claim" from "not yet run live" to the actual result, including the one real defect the run surfaced and fixed along the way — framed as evidence the chain can work, not a reliability guarantee.

Version	Date	Changes
4.15	2026-09-13	B1 (first live end-to-end run, a downstream project, device-management-ui): found a real Mandatory-checkpoints gap. <code>/grill-with-docs</code> wrote a hand-authored Blocked by: (none — API exists, see spec.md) on an unblocked ticket — check 6's own <code>grep</code> passed (the line parses), but <code>adw_watch.py</code>'s <code>BLOCKED_BY_RE</code> splits its value on commas into fragments and treats each as a real blocker; neither fragment matched a sibling ticket, so the ticket sat at <code>ready-for-agent</code>, passed every checkpoint, and was still silently invisible to the frontier scan (<code>just watch --once</code> reported “queue empty” with no error). Added a “Check 6's sharp edge” callout: an unblocked ticket must omit the Blocked by: line entirely, never write “none” in prose. Also confirmed live: <code>/grill-with-docs</code> can write <code>spec.md</code> and the ticket itself inline during grilling, rather than stopping for separate <code>/to-spec /to-tickets</code> runs — “The whole chain, named once” describes 4 distinct human-run steps, but a single grilling session can legitimately collapse steps 1–3 into one. Not yet reflected in that list; flagging here pending a decision on whether to document it as a valid shortcut or leave the list as the general case.

Version	Date	Changes
4.14	2026-09-13	Group A of a fourth-pass re-audit's sync plan (<code>plans/pocock-protocol-sssf-integration.md</code> , <code>downloads/pocock-sssf-sync-plan-v2.md</code>), per explicit direction that the playbook carry no historical naming baggage as live instruction: deleted the "Skill-name compatibility" section outright (the migration it bridged is complete everywhere else in the system; the table's own standing vendoring instruction had been inverted since R2 — history stays only in this changelog's v4.6 row). Fixed the one cross-reference that pointed at it. Rewrote Parts A4/A5/B1's walkthroughs from prose to numbered command steps naming only current skills (<code>to-spec</code> / <code>to-tickets</code> , not the retired <code>write-a-prd</code> / <code>prd-to-plan</code> , which R6 deleted from the local install entirely). Added a new "The whole chain, named once" list at the top of Part D — the working <code>grill-with-docs</code> → <code>to-spec</code> → <code>to-tickets</code> → <code>triage</code> → <code>watch</code> sequence previously had to be assembled from three separate sections. Swept the remaining live retired-name references (Problem 1's opening, the Mandatory-checkpoints trigger line, the queue-gap paragraph, the definition-of-done). Also fixed <code>references/config.md</code> , which still claimed only 3 coding agents exist and that <code>skill_engineering</code> / <code>harness_engineering</code> are ignored under <code>pi</code> / <code>agy</code> — the shipped code covers all four (<code>agent_opencode.py</code> was undocumented entirely); added a full <code>opencode</code> subsection matching <code>agy</code> 's detail level.

Version	Date	Changes
4.13	2026-09-13	F4 (same re-audit): Problem 1's AFK bullet list taught only the interview-related overrides (<code>wayfinder no-fog</code> , <code>to-spec never-prompts</code> , <code>to-tickets quiz-skip</code>), omitting the two triage-bypass overrides R2 actually shipped in the real template (<code>force needs-triage</code> , <code>force plain Status: / Blocked by:</code> lines). An adopter following this list to hand-write their own <code>system.md</code> would get a queue-bypassing, queue-invisible configuration even though the shipped template is correct. Added both missing overrides so the playbook teaches what the template actually does.
4.12	2026-09-13	F2 (same re-audit): fixed a third instance of the exact self-inflicted-error pattern this effort keeps catching — Part D's "judgment call stays interactive" section claimed " <code>to-spec</code> 's interview had no one to answer it headless," which v4.9's mechanical rename introduced and v4.10's own cleanup pass missed; <code>to-spec</code> never interviews, directly contradicting the compatibility table 280 lines earlier. Reworded to point at <code>wayfinder</code> 's fallback instead, which is the thing that's actually still true. Also fixed two wording nits the re-audit flagged in the same pass: the compatibility table's "renamed two of the three" (only one skill was renamed; the other has no upstream equivalent at all) and Problem 1's garbled <code>to-spec / write-a-prd</code> parenthetical.

Version	Date	Changes
4.11	2026-09-13	F1 (independent re-audit of R1–R7, <code>downloads/pocock-sssf-reaudit.md</code>): documented a gap R2’s headless-only planner override left open — <code>/to-spec / /to-tickets</code> self-apply <code>ready-for-agent</code> by default, and running either interactively (which the Filing section’s own guidance recommends) bypasses <code>/triage</code> entirely. Added a “sharp edge” paragraph to Filing: <code>/to-tickets</code>’ bold <code>**Status:** / **Blocked by:**</code> template makes an interactively-filed ticket invisible to <code>adw_watch.py</code>’s frontier scan (a silent stall); plain lines instead produce an untriaged <code>ready-for-agent</code> ticket the watcher will build unjudged. Instructs running <code>/triage</code> on interactive output, or hand-fixing the status line and format, before <code>just watch</code> sees it.
4.10	2026-09-13	R7: converged remaining “PRD” prose on “spec” (upstream finished this same rename in its 1.2.0 — the bootstrap/AFK audit’s “one artifact, four names” finding was the same disease). Problem 1’s title and body still used the retired <code>write-a-prd / prd-to-plan</code> names throughout (R1 deliberately left prose like this alone in favor of a compatibility table) — since a PRD→spec wording pass sitting right next to unrenamed skill names would read incoherently, renamed those too in this one section: <code>write-a-prd</code> → <code>to-spec</code>, <code>prd-to-plan</code>’s “ask the user to paste it” → <code>to-tickets</code>’s “quiz the user”. Scoped to Problem 1 and one Filing-section mention only, not a playbook-wide sweep — the ~15 remaining <code>/write-a-prd / /prd-to-plan</code> mentions elsewhere (mostly diagrams and Parts A/B walkthroughs) stay as-is per R1’s original “documentation-only, point to the compatibility table” decision.

Version	Date	Changes
4.9	2026-09-13	R5: updated the local <code>trriage</code> skill install (<code>~/agents/skills/trriage</code> — the real location; <code>~/claude/skills/trriage</code> is a symlink to it) to the fresh upstream clone, which adds external-PR-as-request-surface support (<code>disable-model-invocation: true</code> , off by default). Backed up the prior version alongside it before overwriting. Added a Part D paragraph noting the PR surface exists but hits a real ceiling today: a PR only exists on a GitHub/GitLab tracker, and <code>adw_watch.py</code> currently refuses those outright (its own module docstring already flags GitHub/GitLab support as unbuilt) — <code>trriage</code> can flip a PR to <code>ready-for-agent</code> , but nothing in this repo's queue would ever pick it up yet. Also fixed one more stale <code>write-a-prd</code> reference (the queue's “judgment call stays interactive” section) to <code>to-spec</code> .
4.8	2026-09-13	Formally accepted the <code>/grill-with-docs</code> → <code>/trriage</code> handoff as manual (R4 in <code>plans/pocock-protocol-sssf-integration.md</code> , citing the bootstrap/AFK audit's finding 3.1.3) rather than building a bridge skill — evaluated and rejected as not worth maintaining for one reminder a checklist line covers as well. Added a new subsection in Part C recording the decision (checked against current upstream Pocock skills too: nothing there bridges it either), and a new Definition-of-done checklist item. Also fixed another stale <code>write-a-prd</code> / <code>prd-to-plan</code> reference (Problem 2's composition-list note) missed by the v4.6 rename pass, to <code>to-spec</code> / <code>to-tickets</code> .

Version	Date	Changes
4.7	2026-09-13	Added a “Two owners of one file” subsection to the Filing section: <code>docs/agents/issue-tracker.md</code>’s base belongs to <code>/setup-matt-pocock-skills</code>, SSSF’s queue depends on that same file also carrying its own extensions (the <code>claimed / resolved states adw_watch.py</code> writes), and nothing merges the two automatically — a re-sync silently drops the extension. Documents the <pre><!-- sssf:queue-extension --> marker</pre> convention (same idiom as <code>vendor_skill.py</code>’s provenance headers) as a diff aid, not automatic reinjection. Also fixed a stale <code>/write-a-prd</code> reference in the Filing section’s shape-picking guidance to <code>/to-spec</code> (missed by the v4.6 rename pass).
4.6	2026-09-13	Added a “Skill-name compatibility” subsection to Part C, right before Problem 1: Matt Pocock’s upstream skill collection has renamed <code>write-a-prd</code> to <code>to-spec</code> (via an intermediate <code>to-prd</code>) and has no equivalent for <code>prd-to-plan</code> (hand-authored here; nearest upstream behavior is <code>to-tickets + implement</code>). Verified against a fresh clone of the upstream repo plus its CHANGELOG, cross-checked in <code>.okf/pocock-skills/naming-drift.md</code>. This playbook’s Part C prose and the planner’s “On the skills composed below” section were both written against the old names — a new adopter vendoring from current upstream couldn’t find two of the three skills this Part tells them to. Documentation-only; the vendored files and planner prompt in this repo are unchanged (<code>prd-to-plan</code> stays load-bearing until a deliberate follow-up migration).

Version	Date	Changes
4.5	2026-09-10	Corrected v4.2's Filing claim: <code>/trriage</code> does write <code>issues/NN-slug.md</code> itself — verified against a real repo's filed tickets (a downstream project), whose content included a <code>.out-of-scope/</code> prior-rejection check and the exact category/state role vocabulary that only <code>/trriage</code>'s documented flow produces, never touched by any other skill. <code>docs/agents/issue-tracker.md</code>'s own “publish to the issue tracker → create a new file” rule means <code>/trriage</code> posting its agent brief on a not-yet-tracked item <i>is</i> the file-creation event on the local-markdown tracker. The real gap is narrower and upstream: nothing hands <code>/trriage</code> the settled description in the first place after <code>grill-with-docs / improve-codebase-architecture</code> finish. Updated the Filing section and Part D's diagram node accordingly. Added a “Mandatory checkpoints” subsection: concrete shell commands to verify a filing actually landed (directory exists, <code>Status:</code> lines present and canonical, nothing uncommitted, <code>Blocked by:</code> references resolve) — this class of failure (claimed-but-not-actually-written, or written-but-invisible-to-the-state-machine) is silent, not an error, the same manufactures-confidence risk rule zero exists to catch for quality gates.

Version	Date	Changes
4.4	2026-09-10	<code>just watch</code> now dispatches through the full SDLC chain (<code>adw_simple_sdlc.py</code> : planner → builder → reviewer → revision loop → documenter → commit), not the lighter <code>plan_build_test</code> chain it silently used before. Found downstream (a downstream project): its first two real queue dispatches (a PIN-authentication access gate, a biometric-unlock follow-up) both landed with zero review — the watcher had always called <code>adw_plan_build_test.main()</code>, which has no reviewer, revision loop, or documenter phase at all. That's a materially bigger gap for an unattended queue than for a manual <code>just sdlc</code> run a human reads afterward. <code>adw_simple_sdlc.main()</code> 's signature is identical (<code>prompt, config, adw_id</code>), confirmed before switching; the roster already had <code>reviewer / documenter</code> configured, so no config change was needed. Updated Part D's "Dispatch" step description to match.

Version	Date	Changes
4.3	2026-09-09	Renamed the <code>just sssf</code> recipe to <code>just watch</code>, matching its script (<code>adw_watch.py</code>) — found via a downstream adoption (a downstream project) questioning the mismatch: 6 of 8 recipes mirror their script name directly, and this one didn't need to be the exception <code>sdlc</code> legitimately is (that one names the workflow's meaning, not its script). Updated the recipe, <code>adw_watch.py</code>'s own runtime log-line prefixes, <code>test_watch.py</code>, and every prescriptive (non-changelog) reference in this playbook's Part D prose and diagram. Also ported back a real bug fix found downstream: <code>git_helper.py</code>'s <code>_git()</code> did a full <code>.strip()</code> on subprocess output, which silently ate the leading space off only the <i>first</i> line of multi-line porcelain output (git status codes are leading-whitespace-significant) — corrupting <code>changed_files()</code>'s first result. Never reachable before a target repo actually had git history to run these functions against; fixed to <code>.rstrip()</code>.
4.2	2026-09-09	Named the filing step Part D's diagram had glossed over: neither <code>grill-with-docs</code> nor <code>improve-codebase-architecture</code> writes an issue — someone has to manually create <code>.scratch/<feature>/spec.md</code> or <code>issues/NN-slug.md</code> with a parseable <code>Status: needs-triage</code> line before <code>/triage</code> can see it. Added a “Filing” subsection with the exact required shape and a new diagram node.
4.1	2026-09-09	Built Part D's <code>just sssf</code> for real: <code>adws/adw_watch.py</code> (stamped like every other ADW) plus 24 tests covering frontier/blocking/claim/resolve/dispatch/tracker-detection. Updated “What this playbook does not claim” — it's shipped and tested, just not yet run against a real live queue end to end.

Version	Date	Changes
4.0	2026-09-09	Added Part D — the queue: turns Part C's single unattended run into a standing queue. Keeps feasibility/compatibility/compliance/security judgment interactive via the <code>triage</code> skill (terminal state <code>ready-for-agent</code> posts a durable agent brief, same bootstrap-then-headless pattern as <code>grill-with-docs</code>); <code>just sssf</code> only ever claims already- <code>ready-for-agent</code> work. Reuses wayfinder's existing map/child/frontier/claim/resolve mechanism as the queue rather than inventing a new one, and flags the gap it exposes: <code>prd-to-plan</code> phases need their own issue files with <code>Blocked by:</code> to be queue-pickable. Design only, not yet built.
3.1	2026-09-08	Replaced every <code>/grill-me</code> reference (A4, A5, B1) with <code>/grill-with-docs</code> . <code>grill-me</code> is a bare alias for the <code>grilling</code> interview with no artifact output; <code>grill-with-docs</code> composes the same interview with <code>domain-modeling</code> , writing resolved vocabulary to <code>CONTEXT.md / docs/adr/</code> . Part C's AFK mechanism depends on that vocabulary existing on disk for scout/planner to read — <code>grill-me</code> alone can't produce it, so the playbook now names one interview skill throughout, and it's the AFK-sufficient one.

Version	Date	Changes
3.0	2026-09-08	Added Part C — going dark: relocates <code>write-a-prd / wayfinder</code> 's interactive interview to a human-supervised bootstrap phase (<code>/grill-with-docs</code> , writes <code>CONTEXT.md / docs/adr/</code> , committed before the headless loop starts) so the AFK loop never needs to prompt live; adds glossary → knowledge-source → structural-search discovery to scout before planning, so plans scope to the delta instead of re-implementing what exists. Generalized from a same-session design pass on a real project (a downstream project's <code>docs/agents/dark-factory-protocol.md</code>), which stays as that project's concrete instance of this part.
2.1	2026-08-31	Removed the last cost figures so the measure-it-yourself stance is consistent. Reframed gate predictions as conditional instructions. Propagated the grill → spec → slice order into Part A as a new A4 step. Unified naming on <code>/grill-me</code> , <code>/write-a-prd</code> , <code>/prd-to-plan</code> and on “slice” as the work unit. Moved rule zero early in the entry diagram to match the prose. Added verification pointers for every mechanism claim, a worked grilled-vs-ungrilled spec example, and made the <code>--setting-sources</code> note self-contained.
2.0	2026-08-31	Rewrote as a general adoption playbook: removed session- and repo-specific anecdotes and cost figures in favour of measure-it-yourself guidance. Corrected the Part B interactive workflow order (grill → spec → plan) and each step's purpose. Rewrote “Where the two layers sit” to describe <code>skill_engineering</code> as shipped — vendoring command, <code>skill_engineering:</code> config key, <code>claude_code</code> -only, per-run cost visibility.

Version	Date	Changes
1.1	2026-08-25	Added five Mermaid diagrams: the entry fork, the gate-verification loop, the four-jobs decision tree, the interactive-vs-headless split, and where the two layers sit today versus after skill_engineering.
1.0	2026-08-25	Initial playbook: existing-codebase adoption, new-project bootstrap, and the standing checklist.